Daily arXiv Report: Sharp Threshold for Universality of Cokernels of Classical Random Matrix Models over the p-adic Integers

Report on `arXiv:2603.12879` by J. Jung, J. Lee, and M. Yu

March 17, 2026

1 Core Question

Given a random $n \times n$ matrix $A(n)$ over the p-adic integers $\mathbb{Z}_p$ whose entries are only mildly random (specifically, α_n-balanced with $\alpha_n = c \log n/n$), what is the *sharp threshold* for the parameter c above which the cokernel $\operatorname{cok}(A(n))$ converges in distribution to the universal limit dictated solely by the symmetry type (non-symmetric, symmetric, or alternating)?

2 Key Definitions and Setup

Definition 1 (p-adic integers and related objects)**.** *Fix a prime p. Let $\mathbb{Z}_p$ denote the ring of p-adic integers, $\mathbb{F}_p = \mathbb{Z}/p\mathbb{Z}$ the finite field with p elements, and $R = \mathbb{Z}/p^d\mathbb{Z}$ for a positive integer d.*

Definition 2 (α-balanced random variable)**.** *A random element $x \in \mathbb{F}_p$ is* α-balanced *(for $\alpha \in (0, 1/2]$) if*

$$\Pr(x = r) \leq 1 - \alpha \quad \textit{for every } r \in \mathbb{F}_p.$$

An element $x \in \mathbb{Z}_p$ is α-balanced if its reduction modulo p is α-balanced. A random matrix in $\mathrm{M}_n(\mathbb{Z}_p)$ (resp. $\mathrm{Sym}_n(\mathbb{Z}_p)$, $\mathrm{Alt}_n(\mathbb{Z}_p)$) is α-balanced if its entries (resp. upper-triangular entries, resp. strictly upper-triangular entries) are independent and α-balanced.

Definition 3 (Cokernel)**.** *For a matrix $A \in \mathrm{M}_n(\mathbb{Z}_p)$ (viewed as a linear map $\mathbb{Z}_p^n \to \mathbb{Z}_p^n$), the* cokernel *is*

$$\operatorname{cok}(A) := \mathbb{Z}_p^n / A(\mathbb{Z}_p^n),$$

which is a finitely generated $\mathbb{Z}_p$-module (and generically a finite abelian p-group).

Intuition. The parameter α controls how "spread out" each entry's distribution is modulo p. When α is a fixed constant, universality of the cokernel distribution was established by Wood and Nguyen–Wood in prior work. The natural question is: how small can $\alpha = \alpha_n$ be (as a function of n) while still having universality? The answer turns out to be $\alpha_n = c \log n/n$ with the critical value at $c = 1$.

3 Early Concrete Example

Example 4 (Non-symmetric case over $\mathbb{F}_p$)**.** *Let $A(n)$ be an α_n-balanced random matrix in $\mathrm{M}_n(\mathbb{F}_p)$ with $\alpha_n = c \log n/n$ and $c > 1$. Then for every nonnegative integer k,*

$$\lim_{n\to\infty} \Pr(\operatorname{rank}(A(n)) = n - k) = \frac{p^{-k^2} \prod_{i=1}^{\infty}(1 - p^{-i})}{\prod_{i=1}^{k}(1 - p^{-i})^2}.$$

This is precisely the rank distribution of a uniformly random *matrix over* $\mathbb{F}_p$*, even though each entry of* $A(n)$ *is allowed to be heavily biased (e.g., taking the value* 0 *with probability* $1 - c\log n/n$*). The result fails at* $c = 1$*.*

4 Main Results

The main theorem treats all three symmetry types simultaneously:

Theorem 5 (Main result, Theorem 1.5)**.** *Let* $c > 1$*,* $\alpha_n = c\log n/n$*, and let* H *be a finite abelian* p*-group. Suppose* $A(n)$ *is an* α_n*-balanced random matrix in* $\mathcal{M}_n(\mathbb{Z}_p)$ *where* $\mathcal{M}_n$ *is one of* M_n*,* Sym_n*, or* Alt_n*.*

- ***Non-symmetric:*** *If* $\mathcal{M}_n = \mathrm{M}_n$*, then*

$$\lim_{n\to\infty} \Pr(\mathrm{cok}(A(n)) \cong H) = \frac{1}{|\mathrm{Aut}(H)|}\prod_{i=1}^{\infty}(1 - p^{-i}).$$

- ***Symmetric:*** *If* $\mathcal{M}_n = \mathrm{Sym}_n$*, then*

$$\lim_{n\to\infty} \Pr(\mathrm{cok}(A(n)) \cong H) = \frac{\#\{\textit{symmetric, bilinear, perfect } \varphi\colon H \times H \to \mathbb{C}^*\}}{|H|\,|\mathrm{Aut}(H)|}\prod_{i=1}^{\infty}(1-p^{1-2i}).$$

- ***Alternating:*** *If* $\mathcal{M}_n = \mathrm{Alt}_n$*, then for* $H \in \mathcal{S}_p$ *(i.e.,* $H \cong G \times G$ *for some finite abelian* p*-group* G*),*

$$\lim_{n\to\infty} \Pr(\mathrm{cok}(A(2n)) \cong H) = \frac{|H|}{|\mathrm{Sp}(H)|}\prod_{i=1}^{\infty}(1 - p^{1-2i}),$$

  and $\Pr(\mathrm{cok}(A(2n)) \cong H) \to 0$ *if* $H \notin \mathcal{S}_p$*.*

Moreover, universality ***fails*** *at the critical scale* $c = 1$*.*

Corollary 6 (Application to sandpile groups)**.** *Let* $\Gamma(n) \in G(n, \beta_n)$ *be an Erdős–Rényi random graph with* $\beta_n \in [\alpha_n, 1 - \alpha_n]$ *and* $\alpha_n = c\log n/n$*,* $c > 1$*. Then the Sylow* p*-subgroup* $S_{\Gamma(n),p}$ *of the sandpile group satisfies*

$$\lim_{n\to\infty} \Pr(S_{\Gamma(n),p} \cong H) = \frac{\#\{\textit{symmetric, bilinear, perfect } \varphi\colon H \times H \to \mathbb{C}^*\}}{|H|\,|\mathrm{Aut}(H)|}\prod_{i=1}^{\infty}(1 - p^{1-2i}).$$

This extends Wood's result (which required constant β_n*) to the sparse regime near the connectivity threshold* $\beta_n = \log n/n$*.*

5 Proof Sketch

The proof proceeds through the following steps:

Step 1: Moment method. The authors use the *moment method for random groups*: if two sequences of random finitely generated $\mathbb{Z}_p$-modules (X_n) and (Y_n) have matching surjective moments $\lim_{n\to\infty} \mathbb{E}[\#\,\mathrm{Sur}(X_n, G)] = M_G$ for all finite abelian p-groups G, and M_G grows at most polynomially in $|G|$, then (X_n) and (Y_n) have the same limiting distribution. This reduces the problem to computing moments of $\mathrm{cok}(A(n))$.

Step 2: Switch from Sur-moments to Hom-moments. A key methodological innovation: instead of computing $\mathbb{E}[\#\,\mathrm{Sur}(\mathrm{cok}(A(n)), G)]$ directly (as in all prior works), the authors work with

$$\mathbb{E}[\#\,\mathrm{Hom}(\mathrm{cok}(A(n)), G)],$$

which is related to Sur-moments by Möbius inversion on the subgroup lattice. This switch is crucial because Hom-moments admit a Fourier-analytic representation that better captures the interaction between the "test homomorphism" $F\colon \mathbb{Z}_p^n \to G$ and the "character" $C \in \mathrm{Hom}(\mathrm{Hom}(\mathbb{Z}_p^n, G), R)$.

Step 3: Unified Fourier-analytic framework. Using the discrete Fourier transform, the Hom-moment is expressed as

$$\mathbb{E}[\#\,\mathrm{Hom}(\mathrm{cok}(A(n)), G)] = \frac{1}{|G|^n} \sum_{\mathbf{g},\mathbf{h}\in G^n} \prod_{(k,l)} \mathbb{E}[\zeta^{A(n)_{kl}\cdot B(\mathbf{g}_k,\mathbf{g}_l)}],$$

where $\zeta = e^{2\pi i/p^d}$ and B is a bilinear form on G^2 that depends on the symmetry type:

- Non-symmetric: $B(\mathbf{g}_k, \mathbf{g}_l) = g_k \cdot h_l$;
- Symmetric: $B(\mathbf{g}_k, \mathbf{g}_l) = g_k \cdot h_l + g_l \cdot h_k$;
- Alternating: $B(\mathbf{g}_k, \mathbf{g}_l) = g_k \cdot h_l - g_l \cdot h_k$.

This formulation treats all cases uniformly.

Step 4: Main term via isotropic subgroups. The *main term* comes from the "trivial part" $\mathcal{G}_{\mathrm{tr}}$: n-tuples $(\mathbf{g}_1, \ldots, \mathbf{g}_n)$ for which $U(\mathbf{g}_k) = 0$ and $B(\mathbf{g}_k, \mathbf{g}_l) = 0$ for all k, l. These correspond to maximal isotropic subgroups of (G^2, B). The main term evaluates to:

- Non-symmetric: $\sum_{H\le G} 1$;
- Symmetric: $\sum_{H\le G} |\wedge^2 H|$;
- Alternating: $\sum_{H\le G} |\mathrm{Sym}^2 H|$.

These counts are related to the universal distributions via the correspondence between maximal isotropic subgroups and bilinear pairings on subgroups of G.

Step 5: Error term analysis (technical core). The bulk of the proof shows that all non-trivial contributions vanish as $n \to \infty$. The authors classify the "non-trivial" n-tuples into four cases based on the "frequency set" $W = \{\mathbf{g} \in G^2 : \#\{k : \mathbf{g}_k = \mathbf{g}\} \ge \gamma n\}$:

(C1) W is **not isotropic**: then $\Theta(n^2)$ pairs (k, l) have $B(\mathbf{g}_k, \mathbf{g}_l) \neq 0$, giving exponential decay $\exp(-\Theta(n \log n))$ via Bauer's maximum principle and the bound $|1 - \alpha_n + \alpha_n \zeta^r| \le 1 - \Theta(\log n/n)$ for $r \neq 0$.

(C2) W is **isotropic but not maximal**: a combinatorial counting argument (using binary entropy bounds) shows the contribution is $o(1)$ for sufficiently small γ.

(C3a–C3c) W is **maximal isotropic** but $(\mathbf{g}_1, \ldots, \mathbf{g}_n) \notin \mathcal{G}_{\mathrm{tr}}$: elements outside W interact non-trivially with W (by the non-constant affine map property of B), producing enough "damping factors" to kill the contribution. Sub-cases handle different regimes of how many elements lie outside W.

(C3d) All $\mathbf{g}_k \in W$ but $(\mathbf{g}_1, \ldots, \mathbf{g}_n) \notin \mathcal{G}_{\mathrm{tr}}$: the tuple doesn't span W, and this is handled by showing the number of such tuples is negligible.

Step 6: Sharpness at $c = 1$. The authors construct explicit α_n-balanced distributions (with $\alpha_n = \log n/n$) for which the cokernel distribution does *not* converge to the universal limit. The key idea is that at $c = 1$, the probability of a zero column (or row) is $\Theta(1/n)$, so the expected number of zero columns remains bounded, preventing universality.

6 Significance

- Resolves open problems of Wood on sharp thresholds for cokernel universality.
- Provides a **unified framework** for non-symmetric, symmetric, and alternating matrices, addressing Wood's Open Problem 3.11.
- The Hom-moment approach (vs. Sur-moments) is a methodological advance that may find further applications.
- The sandpile group application extends the theory to sparse Erdős–Rényi graphs near the connectivity threshold, linking random matrix universality to random graph theory.

Paper: `https://arxiv.org/abs/2603.12879`